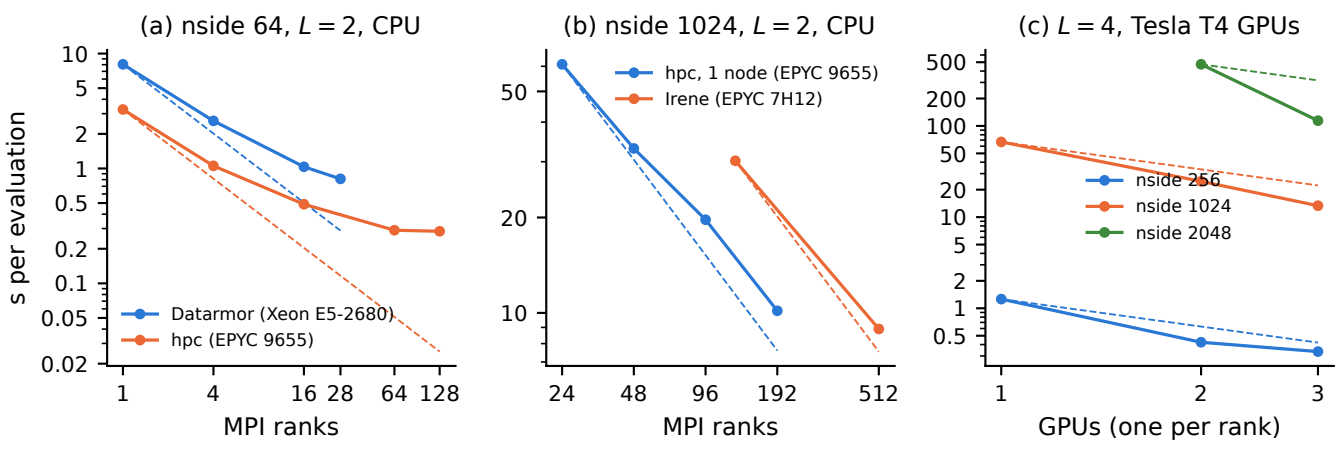
(a) nside 64, $L = 2$, CPU
s per evaluation
10
5
2
1
0.5
0.2
0.1
0.05
0.02
Datarmor (Xeon E5-2680)
hpc (EPYC 9655)
1
4
16
28
64
128
MPI ranks
(b) nside 1024, $L = 2$, CPU
50
20
10
hpc, 1 node (EPYC 9655)
Irene (EPYC 7H12)
24
48
96
192
512
MPI ranks
(c) $L = 4$, Tesla T4 GPUs
500
200
100
50
20
10
5
2
1
0.5
nside 256
nside 1024
nside 2048
1
2
3
GPUs (one per rank)